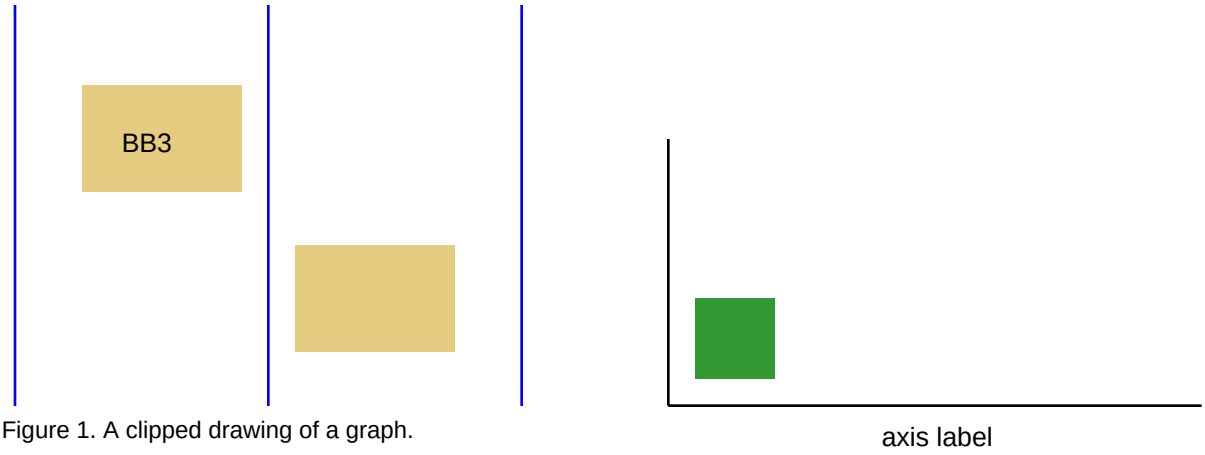

Figure 1. A clipped drawing of a graph.

Once we have the graph with block execution counts and branch probabilities, we make a virtual trace of the memory accesses from it. We use a Python library to read the graph, stored in the dot format. Each node represents a basic block of the target function, and the edges represent the direction from one node to another.